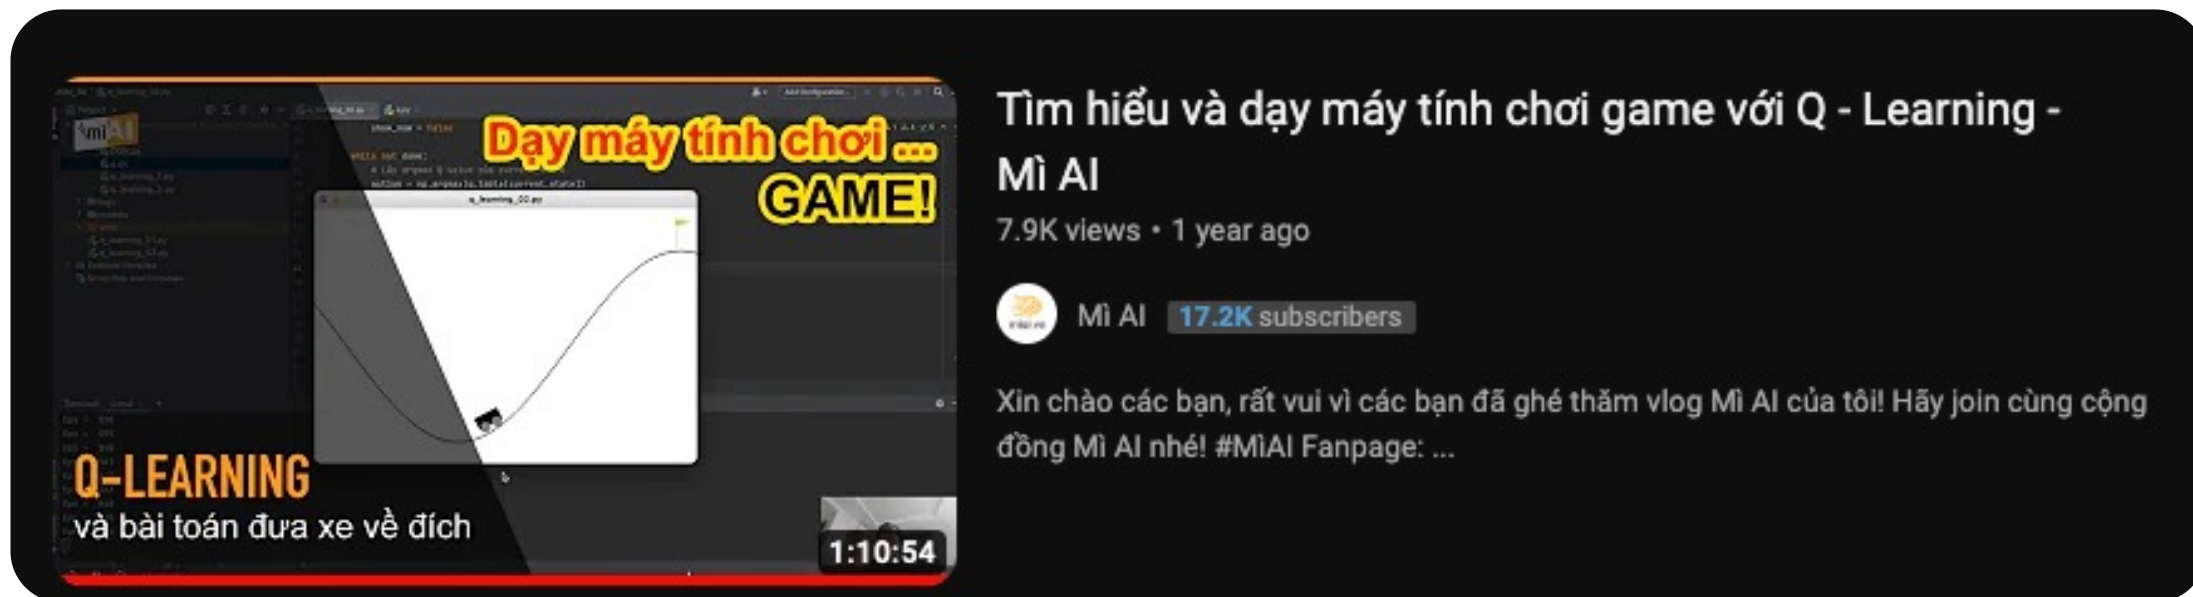




Deep Q-Learning

Mi AI

Q-Learning



The image shows a YouTube video player interface. The video title is "Tìm hiểu và dạy máy tính chơi game với Q - Learning - Mì AI". The video has 7.9K views and was uploaded 1 year ago. The channel is "Mì AI" with 17.2K subscribers. The video description says: "Xin chào các bạn, rất vui vì các bạn đã ghé thăm vlog Mì AI của tôi! Hãy join cùng cộng đồng Mì AI nhé! #MìAI Fanpage: ...". The video thumbnail shows a code editor with the text "Dạy máy tính chơi ... GAME!" and a graph with a curve and a point. The video title in the thumbnail is "Q-LEARNING và bài toán đưa xe về đích". The video duration is 1:10:54.

Q-LEARNING
và bài toán đưa xe về đích

Dạy máy tính chơi ... GAME!

Tìm hiểu và dạy máy tính chơi game với Q - Learning - Mì AI

7.9K views • 1 year ago

Mì AI **17.2K** subscribers

Xin chào các bạn, rất vui vì các bạn đã ghé thăm vlog Mì AI của tôi! Hãy join cùng cộng đồng Mì AI nhé! #MìAI Fanpage: ...

1:10:54


và bài toán đưa xe về đích

Q-LEARNING


1:10:24

channel và video của bạn để...

Q-Learning



Q-Learning



				
1	-	+1	-	+1
2	-	+1	-1	+1
3	-	+1	-1	+1
4	-	+1	-1	-1
5	-	+1	+1	-
...				
23	+1	-	-1	+1
24	+1	-	-1	-1
25	+1	-	+1	-

Q-Learning

		Actions			
		A_1	A_2	...	A_M
States	s_1	$Q(s_1, A_1)$	$Q(s_1, A_2)$		$Q(s_1, A_M)$
	s_2	$Q(s_2, A_1)$	$Q(s_2, A_2)$		$Q(s_2, A_M)$
	\vdots			\ddots	\vdots
	s_N	$Q(s_N, A_1)$	$Q(s_N, A_2)$...	$Q(s_N, A_M)$

Source: <https://en.wikipedia.org/wiki/Q-learning>

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \left(\underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} \right)$$

learned value

Q-Learning

Game Board:



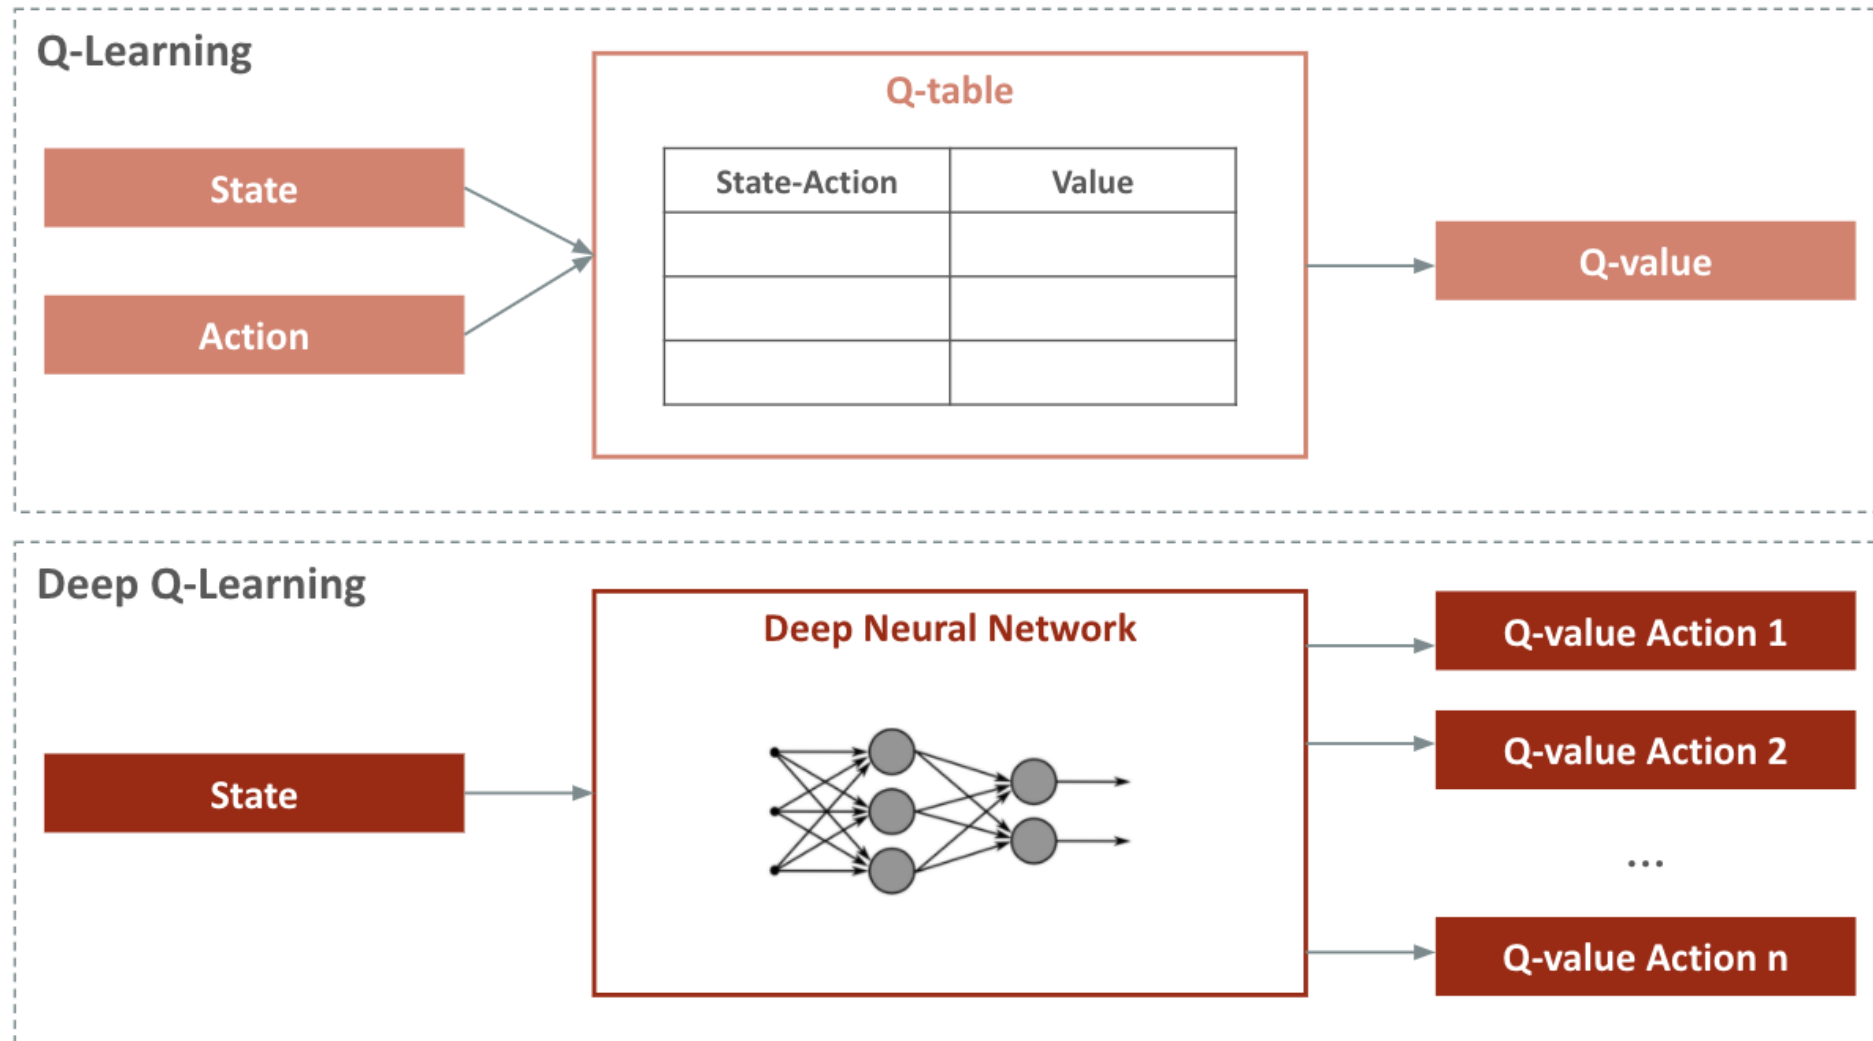
Current state (s):
0 0 0
0 1 0

Q Table:

$\gamma = 0.95$

	0 0 0 1 0 0	0 0 0 0 1 0	0 0 0 0 0 1	1 0 0 0 0 0	0 1 0 0 0 0	0 0 1 0 0 0
↑	0.2	0.3	1.0	-0.22	-0.3	0.0
↓	-0.5	-0.4	-0.2	-0.04	-0.02	0.0
→	0.21	0.4	-0.3	0.5	1.0	0.0
←	-0.6	-0.1	-0.1	-0.31	-0.01	0.0

Deep Q-Learning



Why DQN?

- In an environment with a continuous state space it is impossible to go through all the possible states and actions repeatedly, since there are an infinite number of them and the Q-Table would be too big.
- DQN solves this problem by approximating the Q-Function through a Neural Network and learning from previous training experiences, so that the agent can learn more times from experiences already lived without the need to live them again, as well as avoiding the excessive computational cost of calculating and updating the Q-Table for continuous state spaces.

DQN Components

Main Neural Network

The Main NN tries to predict the expected return of taking each action for the given state.

Train and update every episodes

DQN Components

Target Neural Network

The Target Neural Network is used to get the target value for calculating the loss and optimizing it.

Will be updated every N timesteps with the weights of the main network.

DQN Components

Replay Buffer

The Replay Buffer is a list that is filled with the experiences lived by the agent.

An experience is represented by the **current state**, the **action taken in the current state**, the **reward** obtained after taking that action, whether it is a **terminal state** or not, and the **next state reached** after taking the action.

DQN Components

- State size
- Action size
- Gamma
- Episode
- Number of steps
- Epsilon value, epsilon decay
- Learning rate
- Target NN update rate

Action	Action Number
0	Push cart left
1	Push cart right

Index In array	Meaning	Min Value	Max value
0	Cart Position on x axis	-4.8	4.8
1	Cart Velocity on x axis	$-\infty$	∞
2	Pole Angle	-0.418 rad	0.418 rad
3	Pole Angular Velocity	$-\infty$	∞

DQN Flow

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

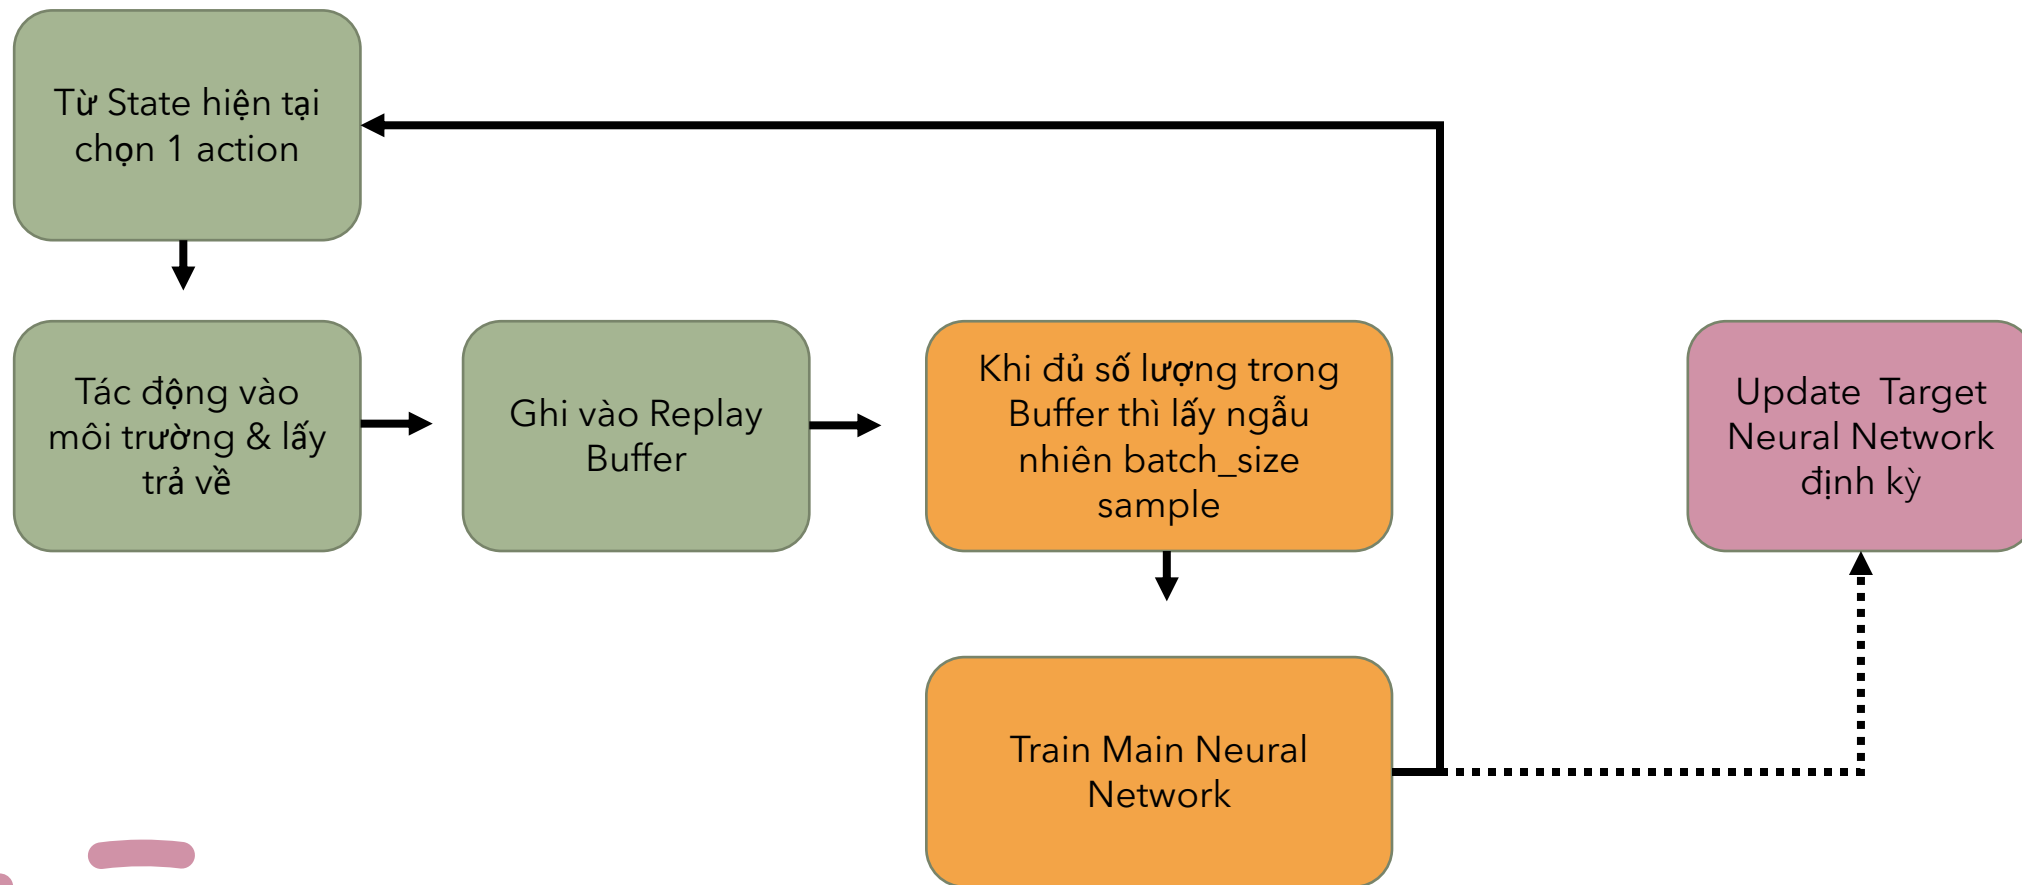
 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for

DQN Flow



Select Action

Epsilon Greedy Policy

1. Generate random number x between 0 and 1
2. $action = \begin{cases} \text{random action} & \text{if } x < \epsilon \\ \text{action with best } Q - \text{Value} & \text{if } x \geq \epsilon \end{cases}$

Train Main NN

$$\text{Loss} = \left(\underbrace{r}_{\text{reward}} + \gamma * \underbrace{\max_{a'} Q(s', a')}_{\substack{\text{Target Network output} \\ \text{for next state and action}}} - \underbrace{Q(s, a)}_{\substack{\text{Main Network output} \\ \text{for current state and action}}} \right)^2$$

$$\text{Set } y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$$

The image features a large white circle centered on a solid orange background. A dashed orange line, composed of several short segments, curves along the top-left edge of the white circle. A solid orange dot is positioned on the right edge of the white circle.

Let's go!